



Get unlimited access

Open in app



Jiahao Meng

Follow

Aug 20 · 8 min read · Listen



Save



# PySpark cheat sheet with code samples

how to initialise Spark, read data, transform it, and build data pipelines In Python.

Markdown Note



Source: [Pyspark](#)

How can we easily parallelize our calculations if our data is too large to work with on a single machine?



115



5





# 1. Introduction

## 1.1 Spark DataFrames VS RDDs

### RDD

Spark's core data structure

✓: A low level object that lets Spark work its magic by splitting data across multiple nodes in the cluster.

✗: However, RDDs are hard to work with directly, so we'll be using the Spark DataFrame abstraction built on top of RDDs.

### Spark DataFrames

Designed to behave a lot like a SQL table

✓:

- easier to understand,
- Operations using DataFrames are automatically optimized
- When using RDDs, it's up to the data scientist to figure out the right way to optimize the query, but the DataFrame implementation has much of this optimization built in!

### Create a SparkSession

- `SparkContext` as our connection to the cluster
- `SparkSession` as our interface with that connection.

# To start working with Spark DataFrames

```
from pyspark.sql import SparkSession
my_spark = SparkSession.builder.getOrCreate()
```





```
print(spark.catalog.listTables())
```

## 2. Spark Schemas

- Define the format of a DataFrame
- May contain various data types:
- Strings, dates, integers, arrays
- Can filter garbage data during import
- Improves read performance

```
# Import the pyspark.sql.types library
from pyspark.sql.types import *

# Define a new schema using the StructType method
people_schema = StructType([
    # Define a StructField for each field
    StructField('name', StringType(), False),
    StructField('age', IntegerType(), False),
    StructField('city', StringType(), False)])

udfpeople = F.udf(peopleParse, ArrayType(people_schema))
```

## 2.1 Transformations

- Lazy
- only executed when we run a Spark action: `.count()`, `.write()`, etc
- can often be modified before being assigned.
- occasionally cause unexpected behaviors:
- IDs not being added until after other transformations have completed
- built-in function: `monotonically_increasing_id()`





```
# new column "Col+1" = Col + 1
df = df.withColumn("Col+1", df.Col + 1)
```

## Filtering

```
long = flights.filter("distance > 1000")
long = flights.filter(flights.distance > 1000)

users_df = users_df.filter(~ col('Name').isNull())
users_df = users_df.filter(users_df.Name.isNotNull())

voter_df = voter_df.withColumn('random_val',
                                when(voter_df.TITLE == 'Councilmember',
                                      F.rand())
                                .when(voter_df.TITLE == 'Mayor', 2)
                                .otherwise(0))
```

## Selecting

- `.select()`
- returns only the columns we specify

```
flights.select("air_time", "origin") flights.select(flights.origin)
flights.select(flights.air_time/60).alias("in_hours")) #SQL expressions
flights.selectExpr("air_time/60 as in_hours")
```

```
#SQL expressions
flights.selectExpr("air_time/60 as in_hours")
```

- `.withColumn()`
- returns all the columns of the DataFrame + we defined

## Aggregating and Grouping

1. `.min()`

2. `.max()`





5. .avg()

# creates a GroupedData to use aboved functions

```
flights.filter(flights.origin == 'PDX').groupBy().min('distance').show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|year|month|day|dep_time|dep_delay|arr_time|arr_delay|carrier|tailnum|flight|o
rigin|dest|air_time|distance|hour|minute|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|2014| 12| 8| 658| -7| 935| -5| VX| N846VA| 1780|
SEA| LAX| 132| 954| 6| 58|
|2014| 1| 22| 1040| 5| 1505| 5| AS| N559AS| 851|
SEA| HNL| 360| 2677| 10| 40|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
```

# Average duration of Delta flights

```
flights.filter(flights.carrier == "DL").groupBy().avg("air_time").show()
```

```
+-----+
| avg(air_time) |
+-----+
|188.20689655172413|
+-----+
```

# Total air\_time in hours: create air\_time/60 column, then sum

```
flights.withColumn("duration_hrs",
```

```
flights.air_time/60).groupBy().sum("duration_hrs").show()
```

```
+-----+
| sum(duration_hrs) |
+-----+
|25289.6000000000126|
+-----+
```

```
by_month_dest = flights.groupBy('month', 'dest')
```

```
by_month_dest.avg('dep_delay').show()
```

```
+-----+-----+-----+-----+
|month|dest| avg(dep_delay)|
+-----+-----+-----+
| 11| TUS| -2.3333333333333335|
| 11| ANC| 7.529411764705882|
| 1| BUR| -1.45|
+-----+-----+-----+-----+
```

```
by_month_dest.agg(F.stddev('dep_delay')).show()
```

```
+-----+-----+-----+-----+
```





## Joining

```
def getFirstAndMiddle(names):  
    # Return a space separated string of names  
    return ' '.join(names[:-1])  
  
# Define the method as a UDF  
udfFirstAndMiddle = F.udf(getFirstAndMiddle, StringType())  
  
# Create a new column using your UDF  
voter_df = voter_df.withColumn('first_and_middle_name',  
  
udfFirstAndMiddle(voter_df.splits))
```

```
+-----+-----+-----+-----+-----+  
+-----+-----+  
|      DATE|      TITLE|      VOTER_NAME|  
splits|first_name|last_name|first_and_middle_name|  
+-----+-----+-----+-----+-----+  
+-----+-----+  
|02/08/2017|Councilmember| Jennifer S. Gates|[Jennifer, S., Gates |  
Jennifer| Gates| Jennifer S.|  
|02/08/2017|Councilmember| Philip T. Kingston|[Philip, T., Kingston |  
Philip| Kingston| Philip T.|  
|02/08/2017| Mayor|Michael S. Rawlings|[Michael, S., Rawlings|  
Michael| Rawlings| Michael S.|  
+-----+-----+-----+-----+-----+  
+-----+-----+  

```

## 2.2 Performance

### Caching



- Stores DataFrames in memory or on disk
- Improves speed on later transformations / actions
- Reduces resource usage





- Local disk based caching may not be a performance improvement
- Cached objects may not be available

## Tips

- Cache only if you need it
- Try caching DataFrames at various points and determine if our performance improves
- Cache in memory and fast SSD / NVMe storage
- Cache to slow local disk if needed
- Use intermediate files!
- Stop caching objects when finished

```
start_time = time.time()

# Add caching
departures_df = departures_df.distinct().cache()

# noting how long the operation takes
print("Counting %d rows took %f seconds" % (departures_df.count(),

time.time() - start_time))
# Counting 139358 rows took 2.679007 seconds

-----

# noting the variance in time of a cached DataFrame
start_time = time.time()
print("Counting %d rows again took %f seconds" % (departures_df.count(),

time.time() - start_time))
# Counting 139358 rows again took 1.184970 seconds

-----

# Determine if is in the cache
departures_df.is_cached)

# Remove from the cache
departures_df.unpersist()
```





- Single node
- Standalone
- Managed
- YARN
- Mesos
- Kubernetes

## 1. Driver Process

- Task assignment
- Result consolidation
- Shared data access
- Tips:
  - Driver node should have double the memory of the worker
  - Fast local storage helpful

## 2. Worker Process

- Runs actual tasks
- Ideally has all code, data, and resources for a given task
- Tips:
  - **More** worker nodes is often better than **larger** workers
  - Test to find the balance
  - Fast local storage extremely useful







- More objects better than larger ones
- Can import via wildcard

```
airport_df=spark.read.csv('airports-*.txt.gz')
```

- General size of objects
- Spark performs better if objects are of similar size

```
# use OS utilities / scripts (split, cut, awk)
split -l 10000 -d largefile chunk-

# to parquet
df_csv = spark.read.csv('singlelargefile.csv')
df_csv.write.parquet('data.parquet')
df = spark.read.parquet('data.parquet')

# read conf setting
spark.conf.get('configuration name')

# write conf setting
spark.conf.set(<configuration name>)
```

## Broadcasting

- Provides a copy of an object to each worker
- Prevents undue / excess communication between nodes
- Can drastically speed up .join() operations

```
from pyspark.sql.functions import broadcast
normal_count = df_1.join(df_2).count()
broadcast_count = df_1.join(broadcast(df_2)).count()

# Normal count:          119910          duration: 3.502130
# Broadcast count:      119910  duration: 1.712519
```





- hides complexity from the user
- Can be slow to complete
- Lowers overall throughput
- Is often necessary, but try to minimize
- LIMIT:
- Limit use of `.repartition (num_partitions)`
- Use `.coalesce (num_partitions)` instead
- Use care when calling `.join()`
- Use `.broadcast()`
- May not need to limit it

## 2.3 Pipelines

### 1. Input(s)

CSV, JSON, web services, databases

```
# Import the data to a DataFrame
departures_df = spark.read.csv('2015-departures.csv.gz', header=True)

# see the column names / order
departures_df.printSchema()

# root
# |-- Date (MM/DD/YYYY): string (nullable = true)
# |-- Flight Number: string (nullable = true)
# |-- Destination Airport: string (nullable = true)
# |-- Actual elapsed time (Minutes): string (nullable = true)

# Remove any duration of 0
departures_df = departures_df.filter(departures_df[3] > 0)

# Add an ID column
departures_df = departures_df.withColumn('id',
F.monotonically_increasing_id())
```





## 2. Transformations

### 2.1 Transformations: `.withcolumn()` , `.filter()` , `.drop()`

- Parse:

#### 1. Incorrect data

- Empty rows
- Commented lines
- Headers

#### 2. Nested structures

- Multiple delimiters:
- 200 300 affenpinscher;0

#### 3. Non-regular data

- Differing numbers of columns per row:
- 600 450 Collie;307 Collie;101 600 449 Japanese\_spaniel;23

#### 4. Focused on CSV data

```
# Can remove comments using an optional argument
df1 = spark.read.csv('datafile.csv.gz', comment='#')

# Handles header fields:
# Defined via argument; Ignored if a schema is defined
df1 = spark.read.csv('datafile.csv.gz', header='True')

# Automatically create columns in a DataFrame based on sep argument df1 =
spark.read.csv('datafile.csv.gz', sep=',')

# Can still successfully parse if sep is not in string
df1 =
spark.read.csv('datafile.csv.gz', sep='*')
```





```
# Split _c0 on the tab character and store the list in a variable
split_cols = F.split(df1['_c0'], '\\t')

# folder, filename, width, height  split_df =
split_cols.withColumn('filename', split_cols.getItem(1))
```

### 3. Validation

```
# Rename the column
df1 = df1.withColumnRenamed('_c0', 'folder')

# Count the number of rows in split_df
split_count = split_df.count()

# Join the DataFrames
joined_df = split_df.join(F.broadcast(df1), folder', 'left_anti')
```

### 4. Output(s)

CSV, Parquet, database

```
# Write the file out to JSON format
departures_df.write.json('output.json', mode='overwrite')
```

### 5. Analysis

- Analysis Calculations (UDF)

```
def getAvgSale(saleslist):
    totalsales = 0
    count = 0
    for sale in saleslist:
        totalsales += sale[2] + sale[3]
    count += 2
    return totalsales / count

udfGetAvgSale = udf(getAvgSale, DoubleType())
df = df.withColumn('avg_sale', udfGetAvgSale(df.sales_list))
```





```
df = df.withColumn('sq_ft', df.width * df.length)
df = df.withColumn('total_avg_size', udfComputeTotal(df.entries) /
df.numEntries)
```

## 3. Pandas

### 3.1 Spark2Pandas

- Spark DataFrames make that easy with the `.toPandas()` method. Calling this method on a Spark DataFrame returns the corresponding pandas DataFrame.

### 3.2 Pandas2Spark

- The `.createDataFrame()` method takes a pandas DataFrame and returns a Spark DataFrame.
- In the last exercise, you saw how to move data from Spark to pandas. However, maybe you want to go the other direction, and put a pandas DataFrame into a Spark cluster! The `SparkSession` class has a method for this as well.
- The `.createDataFrame()` method takes a pandas DataFrame and returns a Spark DataFrame.
- The output of this method is stored locally, not in the `SparkSession` catalog. This means that you can use all the Spark DataFrame methods on it, but you can't access the data in other contexts.
- For example, a SQL query (using the `.sql()` method) that references your DataFrame will throw an error. To access the data in this way, you have to save it as a *temporary table*.
- You can do this using the `.createTempView()` Spark DataFrame method, which takes as its only argument the name of the temporary table you'd like to register. This method registers the DataFrame as a table in the catalog, but as this table is temporary, it can only be accessed from the specific `SparkSession` used to create the Spark DataFrame.
- There is also the method `.createOrReplaceTempView()`. This safely creates a new temporary table if nothing was there before, or updates an existing table if one was already defined. You'll use this method to avoid running into problems with duplicate tables.
- Check out the diagram to see all the different ways your Spark data structures interact with





```
# Create spark_temp from pd_temp
spark_temp = spark.createDataFrame(pd_temp)

# Examine the tables in the catalog
print(spark.catalog.listTables())

# Add spark_temp to the catalog
spark_temp.createOrReplaceTempView("temp")

# Examine the tables in the catalog again
print(spark.catalog.listTables())
```

